

Support Vector Machines: How to compute the margin?

Mining of Massive Datasets
Leskovec, Rajaraman, and Ullman
Stanford University



SVM: How to estimate w ?

$$\min_{w,b} \frac{1}{2} w \cdot w + C \sum_{i=1}^n \xi_i$$

$$s.t. \forall i, y_i \cdot (x_i \cdot w + b) \geq 1 - \xi_i$$

- **Want to estimate w and b !**
 - **Standard way:** Use a solver!
 - **Solver:** software for finding solutions to “common” optimization problems
- **Use a quadratic solver:**
 - Minimize quadratic function
 - Subject to linear constraints
- **Problem:** Solvers are inefficient for big data!

SVM: How to estimate w ?

- **Want to estimate w , b !**

$$\min_{w,b} \frac{1}{2} w \cdot w + C \sum_{i=1}^n \xi_i$$

- **Alternative approach:**

$$s.t. \forall i, y_i \cdot (x_i \cdot w + b) \geq 1 - \xi_i$$


- **Want to minimize $f(w,b)$:**

$$f(w,b) = \frac{1}{2} \sum_{j=1}^d \left(w^{(j)} \right)^2 + C \sum_{i=1}^n \max \left\{ 0, 1 - y_i \left(\sum_{j=1}^d w^{(j)} x_i^{(j)} + b \right) \right\}$$

- **How to minimize convex functions $f(z)$?**

- Use gradient descent: $\min_z f(z)$

- Iterate: $z_{t+1} \leftarrow z_t - \eta f'(z_t)$



SVM: How to estimate w ?

- Want to minimize $f(w, b)$:

$$f(w, b) = \frac{1}{2} \sum_{j=1}^d \left(w^{(j)} \right)^2 + C \sum_{i=1}^n \underbrace{\max \left\{ 0, 1 - y_i \left(\sum_{j=1}^d w^{(j)} x_i^{(j)} + b \right) \right\}}_{\text{Empirical loss } L(x_i, y_i)}$$

- Compute the gradient $\nabla(j)$ w.r.t. $w^{(j)}$

$$\nabla(j) = \frac{\partial f(w, b)}{\partial w^{(j)}} = w^{(j)} + C \sum_{i=1}^n \frac{\partial L(x_i, y_i)}{\partial w^{(j)}}$$

$$\begin{aligned} \frac{\partial L(x_i, y_i)}{\partial w^{(j)}} &= 0 && \text{if } y_i (w \cdot x_i + b) \geq 1 \\ &= -y_i x_i^{(j)} && \text{else} \end{aligned}$$

SVM: How to estimate w ?

■ (Batch) Gradient Descent:

Iterate until convergence:

• **For $j = 1$ to d**

• **Evaluate:** $\nabla(j) = \frac{\partial f(w, b)}{\partial w^{(j)}} = w^j + C \sum_{i=1}^n \frac{\partial L(x_i, y_i)}{\partial w^{(j)}}$

• **Update:**

$$w^{(j)} \leftarrow w^{(j)} - \eta \nabla(j)$$

η learning rate parameter
 C regularization parameter

■ Problem:

■ Computing $\nabla(j)$ takes $O(n)$ time!

■ n ... size of the training dataset

SVM: How to estimate w ?

We just had:

$$\nabla(j) = w^{(j)} + C \sum_{i=1}^n \frac{\partial L(x_i, y_i)}{\partial w^{(j)}}$$

■ Stochastic Gradient Descent

- Instead of evaluating gradient over all examples evaluate it for each **individual** training example

$$\nabla(j, i) = w^{(j)} + C \frac{\partial L(x_i, y_i)}{\partial w^{(j)}}$$



■ Stochastic gradient descent:

Iterate until convergence:

- For $i = 1 \quad n$
 - For $j = 1 \quad d$
 - **Evaluate:** $\nabla(j, i)$
 - **Update:** $w^{(j)} \leftarrow w^{(j)} - \eta \nabla(j, i)$

SGD vs. GD

■ Convergence of **GD** vs. **SGD**



GD improves the value of the objective function at every step.

SGD improves the value but in a “noisy” way.

GD takes fewer steps to converge but each step takes much longer to compute.

In practice, **SGD** is much faster!

Example: Text categorization

- **Example by Leon Bottou:**
 - **Reuters RCV1** document corpus
 - Predict a category of a document
 - One **vs.** the rest classification
 - **$n = 781,000$** training examples (documents)
 - 23,000 test examples
 - **$d = 50,000$** features
 - One feature per word
 - Remove stop-words
 - Remove low frequency words

Example: Text categorization

■ Questions:

- (1) Is **SGD** successful at minimizing $f(\mathbf{w}, \mathbf{b})$?
- (2) How quickly does **SGD** find the min of $f(\mathbf{w}, \mathbf{b})$?
- (3) What is the error on a test set?

	<i>Training time</i>	<i>Value of $f(\mathbf{w}, \mathbf{b})$</i>	<i>Test error</i>
Standard SVM	23,642 secs	0.2275	6.02%
“Fast SVM”	66 secs	0.2278	6.03%
SGD SVM	1.4 secs	0.2275	6.02%

- (1) SGD-SVM is successful at minimizing the value of $f(\mathbf{w}, \mathbf{b})$
- (2) SGD-SVM is super fast
- (3) SGD-SVM test set error is comparable