

PageRank: The Complete Algorithm

- Input: Graph G and parameter β

- Directed graph G with spider traps and dead ends
- Parameter β

- Output: PageRank vector r

- **Set:** $r_j^{(0)} = \frac{1}{N}, \quad t = 1$

- **do:**

- $\forall j: r_j'^{(t)} = \sum_{i \rightarrow j} \beta \frac{r_i^{(t-1)}}{d_i}$

- $r_j'^{(t)} = 0$ if in-deg. of j is 0

- **Now re-insert the leaked PageRank:**

- $\forall j: r_j^{(t)} = r_j'^{(t)} + \frac{1-S}{N}$ **where:** $S = \sum_j r_j'^{(t)}$

- $t = t + 1$

- **while** $\sum_j |r_j^{(t)} - r_j^{(t-1)}| > \varepsilon$

Sparse Matrix Encoding

- **Encode sparse matrix using only nonzero entries**
 - Space proportional roughly to number of links
 - Say $10N$, or $4 \times 10 \times 1$ billion = 40GB
 - **Still won't fit in memory, but will fit on disk**

source node	degree	destination nodes
0	3	1, 5, 7
1	5	17, 64, 113, 117, 245
2	2	13, 23

Basic Algorithm: Update Step

- Assume enough RAM to fit r^{new} into memory

- Store r^{old} and matrix \mathbf{M} on disk


- Then 1 step of power-iteration is:

Initialize all entries of r^{new} to $(1-\beta)/N$


For each page p (of out-degree n):

Read into memory: $p, n, dest_1, \dots, dest_n, r^{old}(p)$

for $j = 1 \dots n$: $r^{new}(dest_j) += \beta r^{old}(p) / n$



src	degree	destination
0	3	1, 5, 6
1	4	17, 64, 113, 117
2	2	13, 23



Analysis

- Assume enough RAM to fit r^{new} into memory
 - Store r^{old} and matrix M on disk
- In each iteration, we have to:
 - Read r^{old} and M
 - Write r^{new} back to disk
 - Input/Output cost = $2|r| + |M|$
- Question:
 - What if we could not even fit r^{new} in memory?

Block-based Update Algorithm

r^{new}	
0	
1	
2	
3	
4	
5	

src	degree	destination
0	4	0, 1, 3, 5
1	2	0, 5
2	2	3, 4

r^{old}	
	0
	1
	2
	3
	4
	5

Analysis of Block Update

- Similar to nested-loop join in databases
 - Break r^{new} into k blocks that fit in memory
 - Scan M and r^{old} once for each block
- k scans of M and r^{old}
 - $k(|M| + |r|) + |r| = k|M| + (k+1)|r|$
- Can we do better?
 - Hint: M is much bigger than r (approx 10-20x), so we must avoid reading it k times per iteration

Block-Stripe Update Algorithm

r^{new}

0	
1	

2	
3	

4	
5	

src	degree	destination
0	4	0, 1
1	3	0
2	2	1

0	4	3
2	2	3

0	4	5
1	3	5
2	2	4

r^{old}

	0
	1
	2
	3
	4
	5

Block-Stripe Analysis

- Break M into stripes
 - Each stripe contains only destination nodes in the corresponding block of r^{new}
- Some additional overhead per stripe
 - But it is usually worth it
- Input/Output cost per iteration
 - $|M|(1+\varepsilon) + (k+1)|r|$

Some Problems with Page Rank

- **Measures generic popularity of a page**
 - Biased against topic-specific authorities
 - **Solution:** Topic-Specific PageRank
- **Uses a single measure of importance**
 - Other models e.g., hubs-and-authorities
 - **Solution:** Hubs-and-Authorities
- **Susceptible to Link spam**
 - Artificial link topographies created in order to boost page rank
 - **Solution:** TrustRank